

# Video Processing

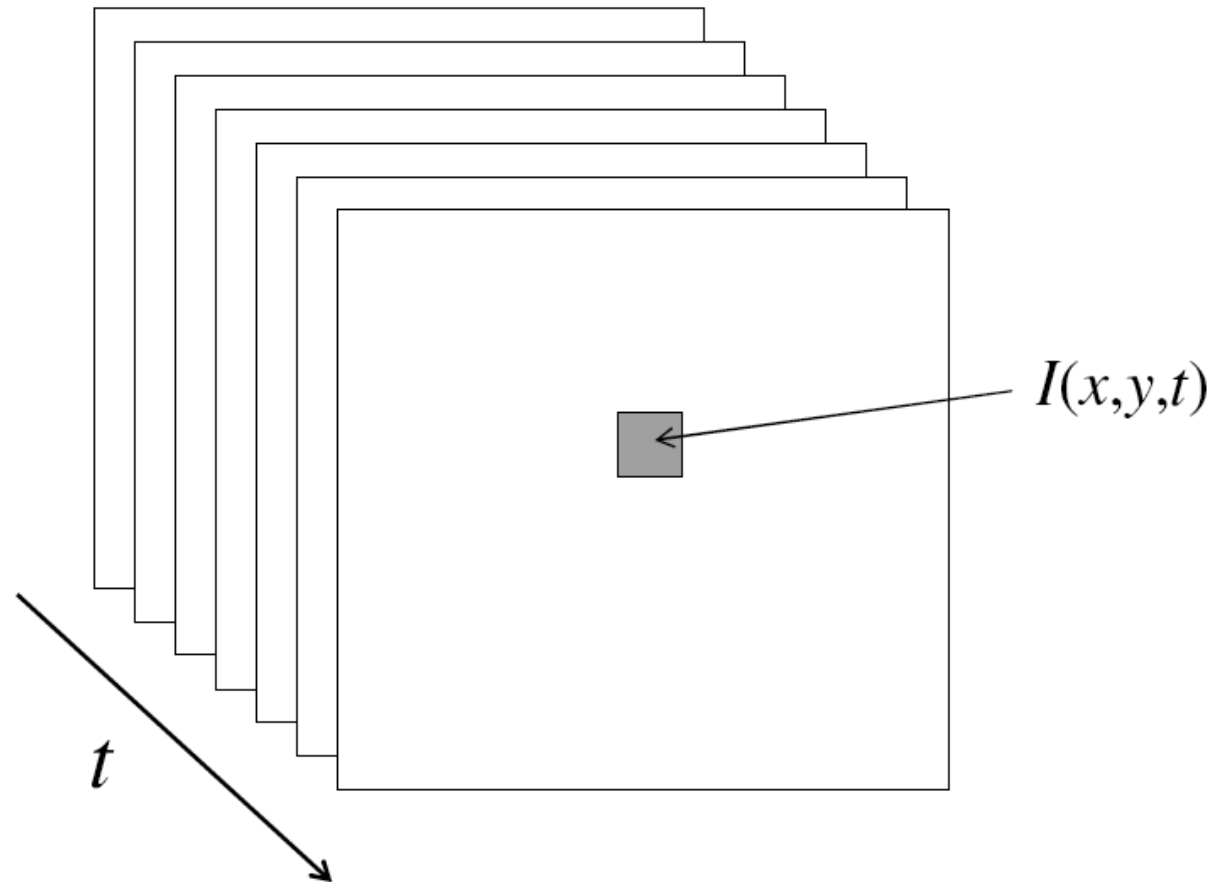
Paulo Dias





- Motion
  - Motion analysis
  - Optical Flow
- Background Subtraction
- Tracking
  - Object Tracking
  - Template Matching
  - Detection vs Tracking

- A video is a sequence of frames **captured over time**
- The image data is a function of **space** (x, y) and **time** (t)





- Several **information** can be extracted from **time varying** sequences of images:
  - **Camouflaged** objects easily visible when moving
  - **Size and position** of objects are more easily determined when the objects move
  - Even **simple** image **differencing** provides edge detector for objects moving over any static background



é a diferença entre os frames - da nos a in

Lucas & Kanade Optical Flow method



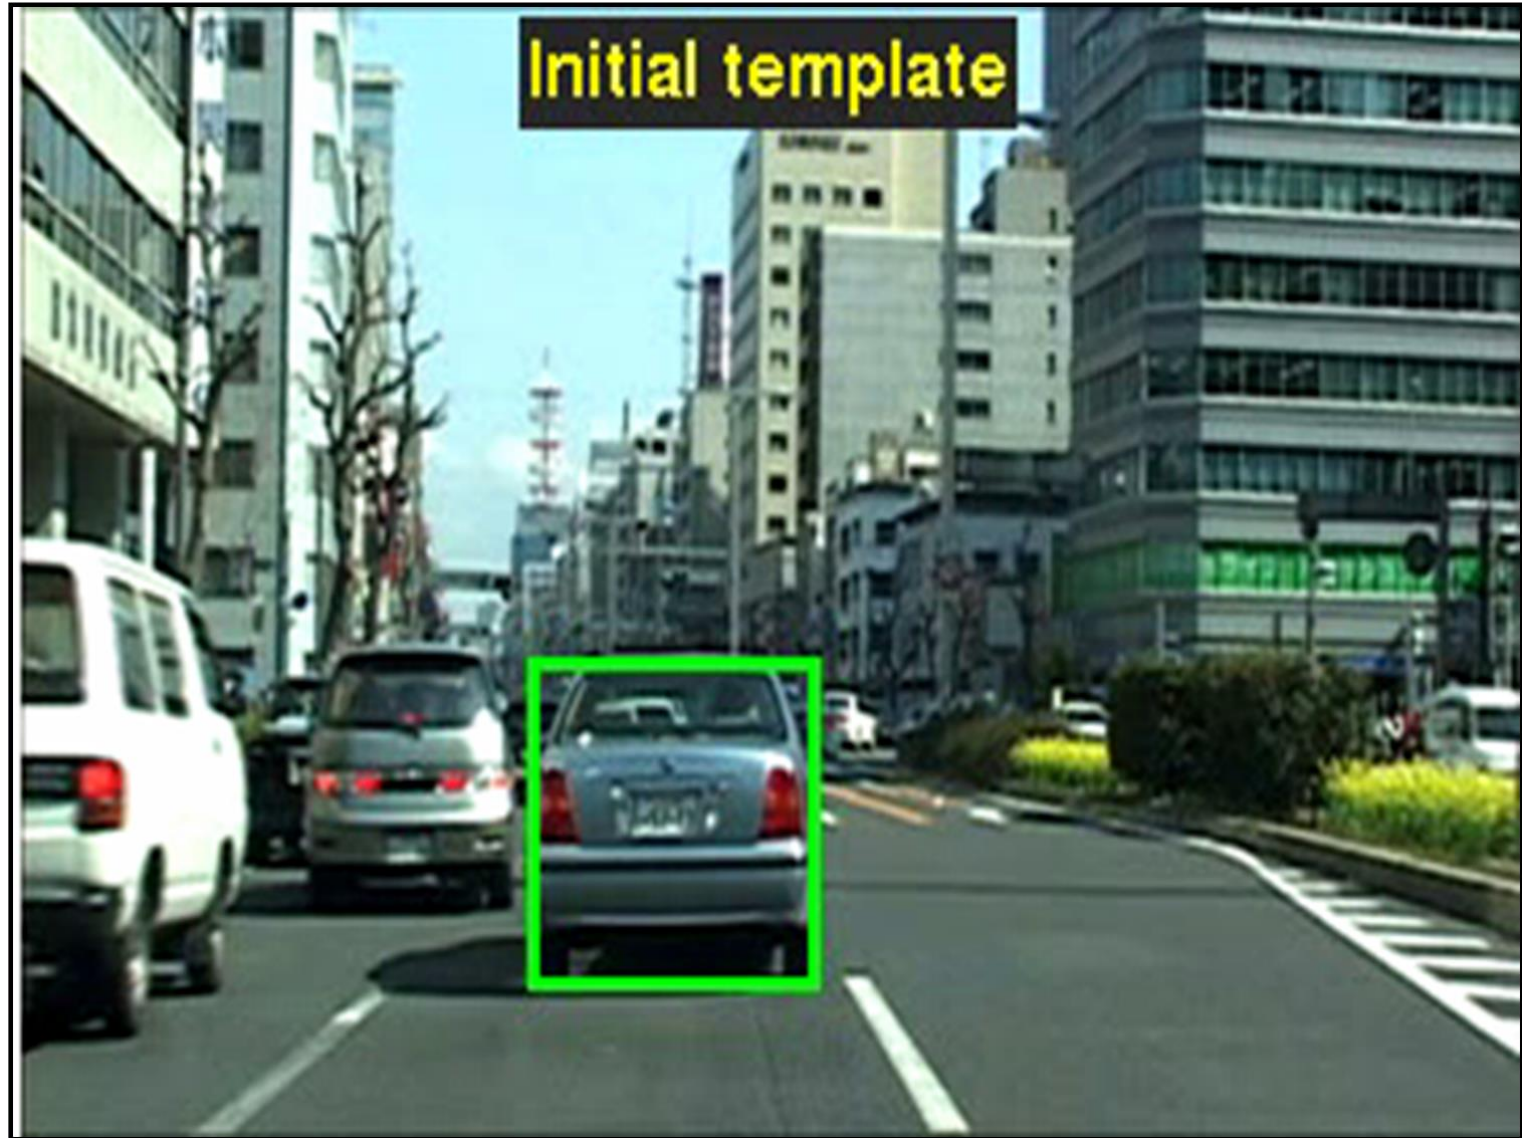
- Analysis of visual motion can be divided into two stages:
  - **measurement** of the motion
  - use of motion data to **segment scene** into objects and to extract information about shape and motion.
- There are two types of motion to consider:
  - movement in the scene = **static** camera,
  - movement of the camera = **ego** motion.
    - should be the same (motion is relative) but not always true - if scene moves illumination, shadow and specular effects need to be considered



- We are interested in **finding** the **movement** of scene objects from time-varying images (videos).
- Lots of uses
  - Track object behavior
  - Align images (mosaics)
  - 3D shape reconstruction
  - Correct for camera jitter (stabilization)
  - Special effects



# Rigid Object Tracking



(Simon Baker, CMU)

# Non Rigid Object Tracking





(Comaniciu et al, Siemens)



[https://youtu.be/Yw\\_zkLaZNSQ](https://youtu.be/Yw_zkLaZNSQ)





# Structure from motion



(David Nister, Kentucky)

- motion field is the **projection** of the 3D scene motion into the image.





- Length of **flow vectors** inversely proportional to **depth  $Z$**  of 3D point - points closer to the camera move more quickly across the image plane.

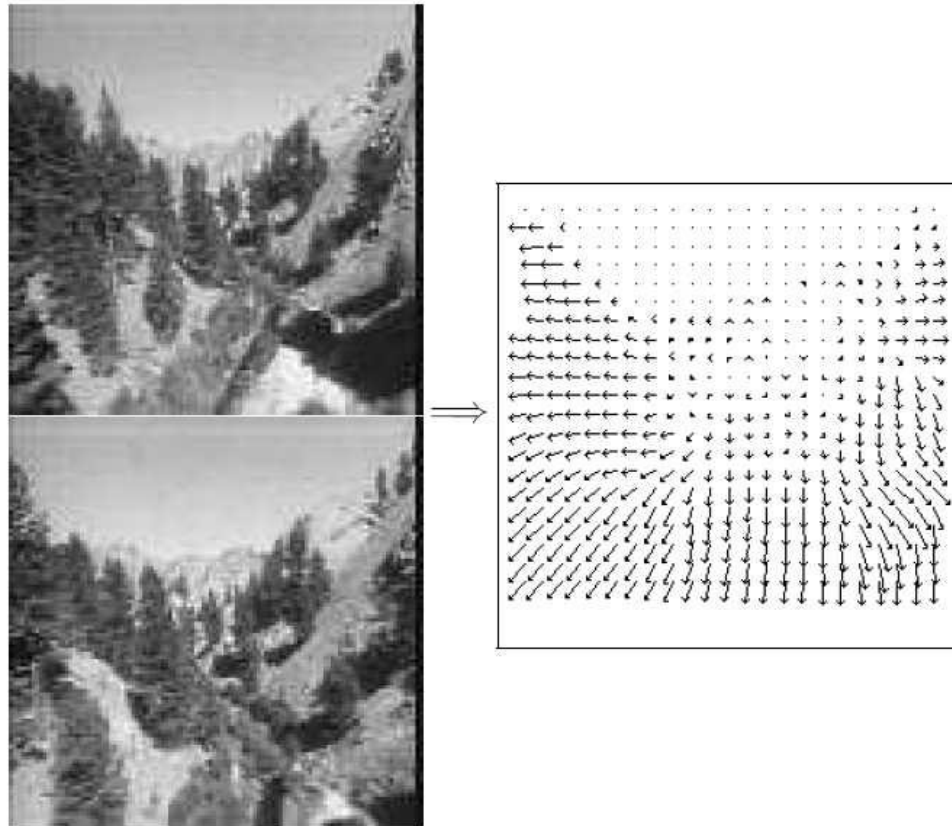


Figure 1.2: Two images taken from a helicopter flying through a canyon and the computed optical flow field.



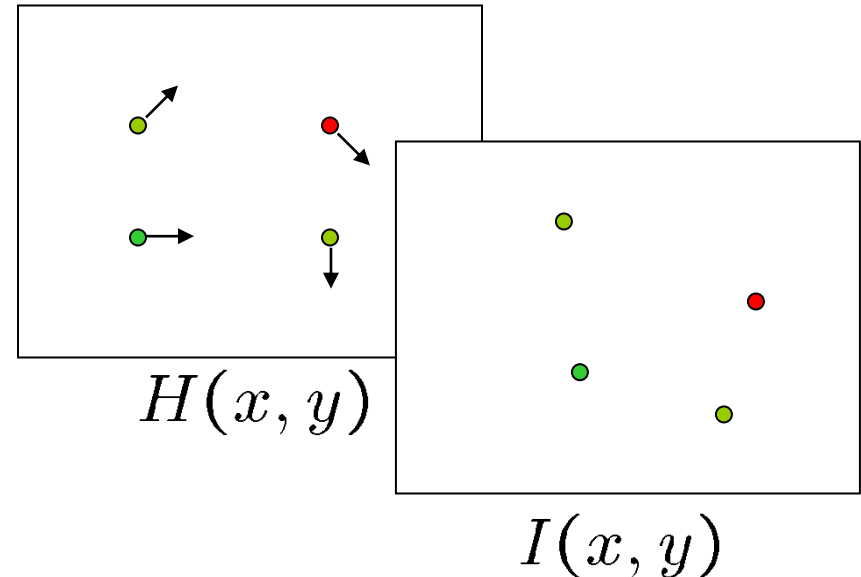
- Optical flow is the apparent motion of brightness patterns (or colours) in the image.
- Ideally, **optical flow** should be the **same** as the **motion field**.
- Careful: apparent motion can be caused by **lighting changes** without any actual motion.
- To estimate pixel motion from image we have to solve the **pixel correspondence** problem.
- Given a pixel in frame  $t$ , **look for nearby pixels** with same characteristics (colour, brightness, ...) in frame  $t - 1$ .



# Problem Definition: Optical Flow



- How to estimate pixel motion from image  $H$  to image  $I$ ?
  - Find pixel correspondences
    - Given a pixel in  $H$ , look for nearby pixels of the same color in  $I$



- Key assumptions
  - **color constancy**: a point in  $H$  looks “the same” in image  $I$ 
    - For grayscale images, this is **brightness constancy**
  - **small motion**: points do not move very far



- How to get more equations for a pixel?
  - Basic idea: impose additional constraints
    - most common is to assume that the flow field is smooth locally
    - one method: pretend the pixel's neighbors have the same  $(u,v)$ 
      - If we use a 5x5 window, that gives us 25 equations per pixel!

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v] \quad \text{color constancy}$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

$\underset{25 \times 2}{A} \quad \underset{2 \times 1}{d} \quad \underset{25 \times 1}{b}$

small motion  
constant  
within  
neighborhood



- Prob: we have more equations than unknowns

$$\begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 \quad 25 \times 1 \end{matrix} \longrightarrow \text{minimize } \|Ad - b\|^2$$

- **Solution: solve least squares problem**

- minimum least squares solution given by solution (in d) of:

$$\begin{matrix} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 \quad 2 \times 1 \end{matrix}$$

$$\begin{matrix} \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} & \begin{bmatrix} u \\ v \end{bmatrix} = - & \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix} \\ A^T A & & A^T b \end{matrix}$$

- The sum are over all pixels in the K x K window
- This technique was first proposed by Lukas & Kanade (1981)
  - described in Trucco & Verri reading

- Optimal  $(u, v)$  satisfies Lucas-Kanade equation

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A$   $A^T b$

## When is This Solvable?

- $A^T A$  should be invertible
- $A^T A$  should not be too small due to noise
  - eigenvalues  $\lambda_1$  and  $\lambda_2$  of  $A^T A$  should not be too small
- $A^T A$  should be well-conditioned
  - $\lambda_1 / \lambda_2$  should not be too large ( $\lambda_1$  = larger eigenvalue)




- Optical flow:
  - Algorithms try to **approximate the true motion field** of the image plane.
  - The Optical Flow Constraint Equation needs **additional constraints** (e.g. smoothness, constant local flow).
  - The **Lucas Kanade** method is the most **popular** Optical Flow Algorithm.
- Lucas Kanade Optical Flow in OpenCV
  - [https://docs.opencv.org/4.x/d4/dee/tutorial\\_optical\\_flow.html](https://docs.opencv.org/4.x/d4/dee/tutorial_optical_flow.html)



- Motion
  - Motion analysis
  - Optical Flow
- Background Subtraction
- Tracking
  - Object Tracking
  - Template Matching
  - Detection vs Tracking

- It is possible to look at video data as a spatio-temporal volume.
- If camera is stationary, each line through time corresponds to a single ray in space.





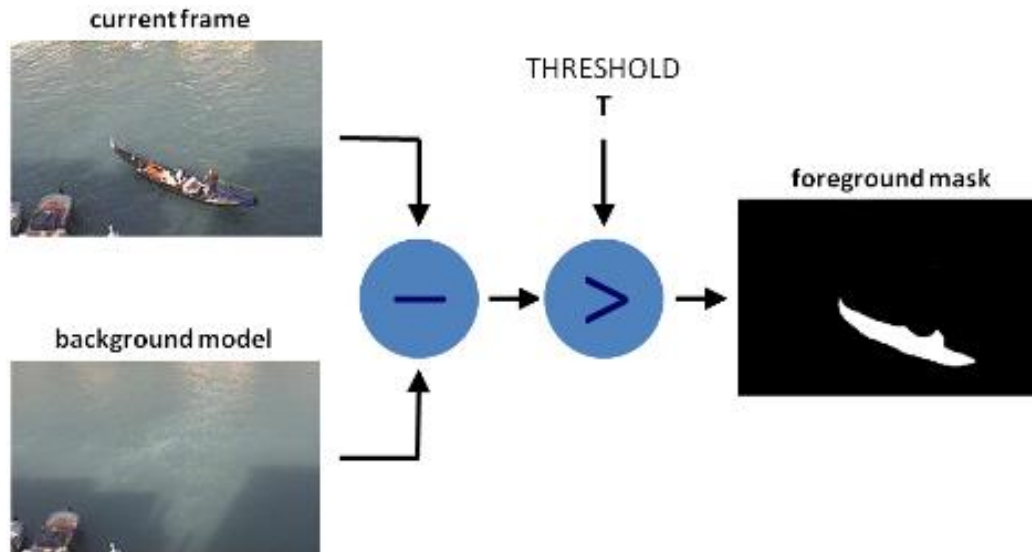
- Background subtraction techniques are commonly used for segmenting out objects of interest in a static camera scene for applications such as:
  - Surveillance
  - Robot vision
  - Object tracking
  - Traffic applications
  - Human motion capture
  - Augmented reality



# Background Subtraction



- Used for generating a **foreground mask** (binary image) containing **moving objects** in static camera setups.
- Name comes from the **simple technique of subtracting** the observed image from the estimated image and thresholding the result to generate the objects of interest.





- Foreground **detection** – how the object areas are distinguished from the background
- Background **maintenance** – how the background is maintained over time
- Post-processing – how the **segmented** object areas are detected

- Generic algorithm:
  - Create an image of the **stationary background** by averaging a long sequence.
  - **Subtract current frame** and known background frame
  - Motion detection algorithms such as these only work if the **camera** is **stationary** and objects are moving against a **fixed background**

Image at time  $t$ :

$$I(x, y, t)$$



Background at time  $t$ :

$$B(x, y, t)$$



—

$$| > Th$$





- Generic algorithm:
  - With **frame differencing**, background is estimated to be the previous frame. Background subtraction equation becomes  $B(x, y, T) = I(x, y, t - 1)$  and  $|I(x, y, t) - I(x, y, t - 1)| > Th$
  - Depending on the object structure, speed, frame rate and global **threshold** may or may not be useful (usually not).
  - Another approach is to **model the background** using a running average. A pixel is marked as foreground if

$$|I_t - B_t| > th$$

where **th** is a predefined threshold.





























The thresholding is often followed by **morphological closing** with a 3x3 kernel and the discarding of small regions

- The **background update** is

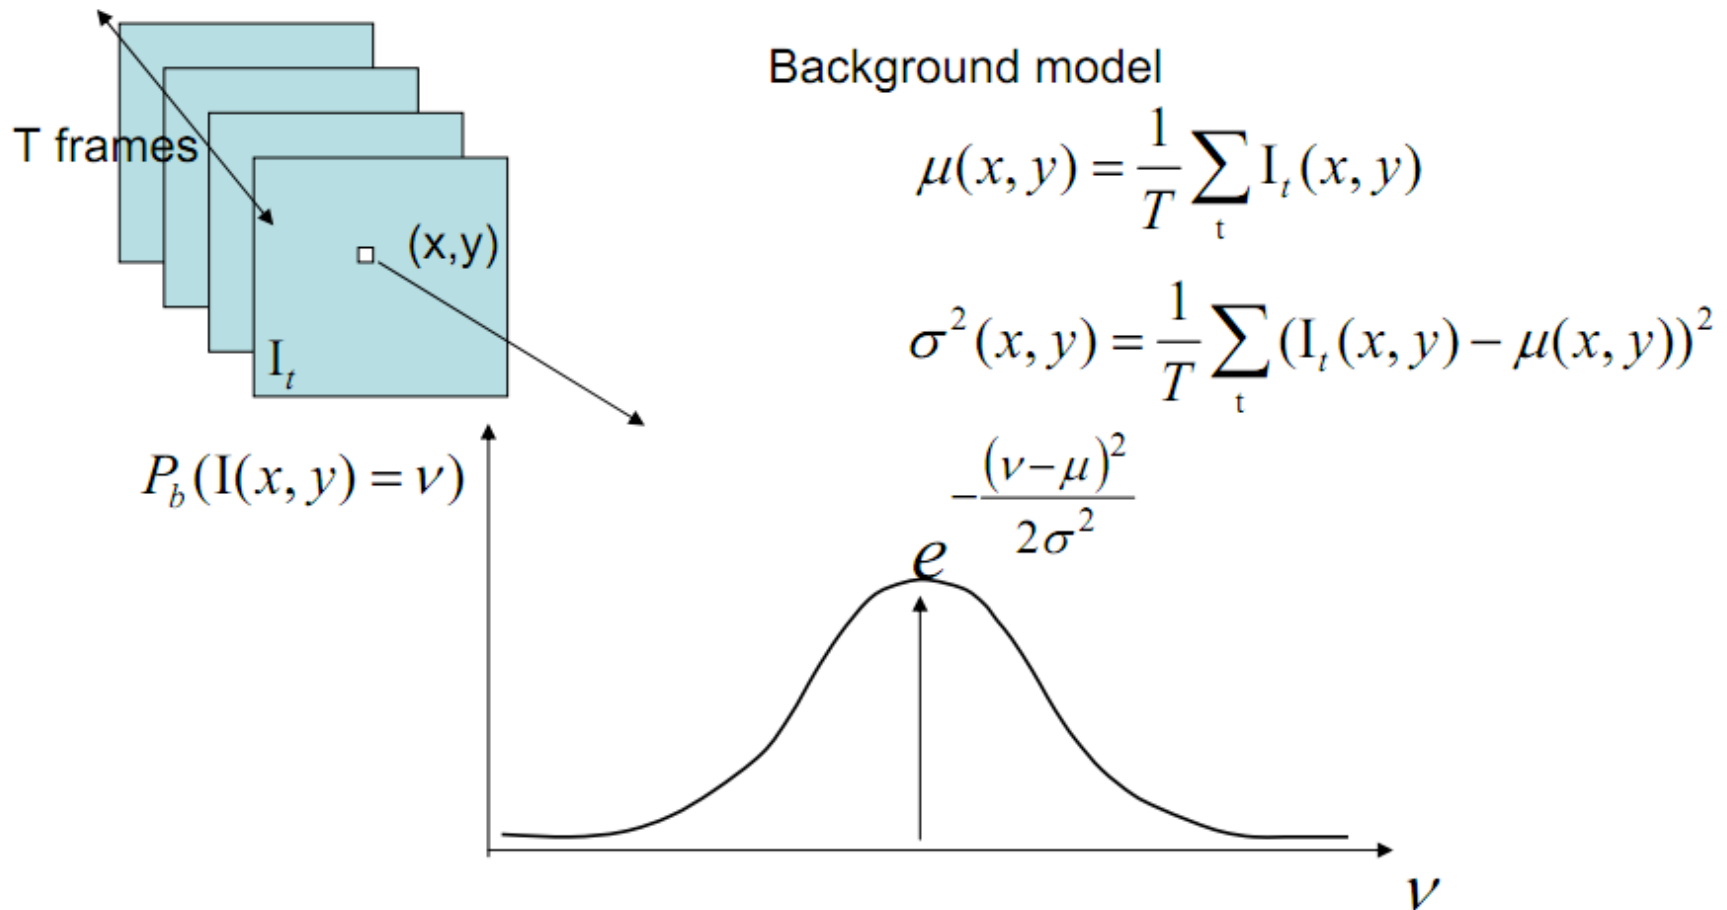
$$B_{t-1} = \alpha I_t + (1 - \alpha) B_t$$

where  $\alpha$  is kept small to prevent artificial tails forming behind moving object

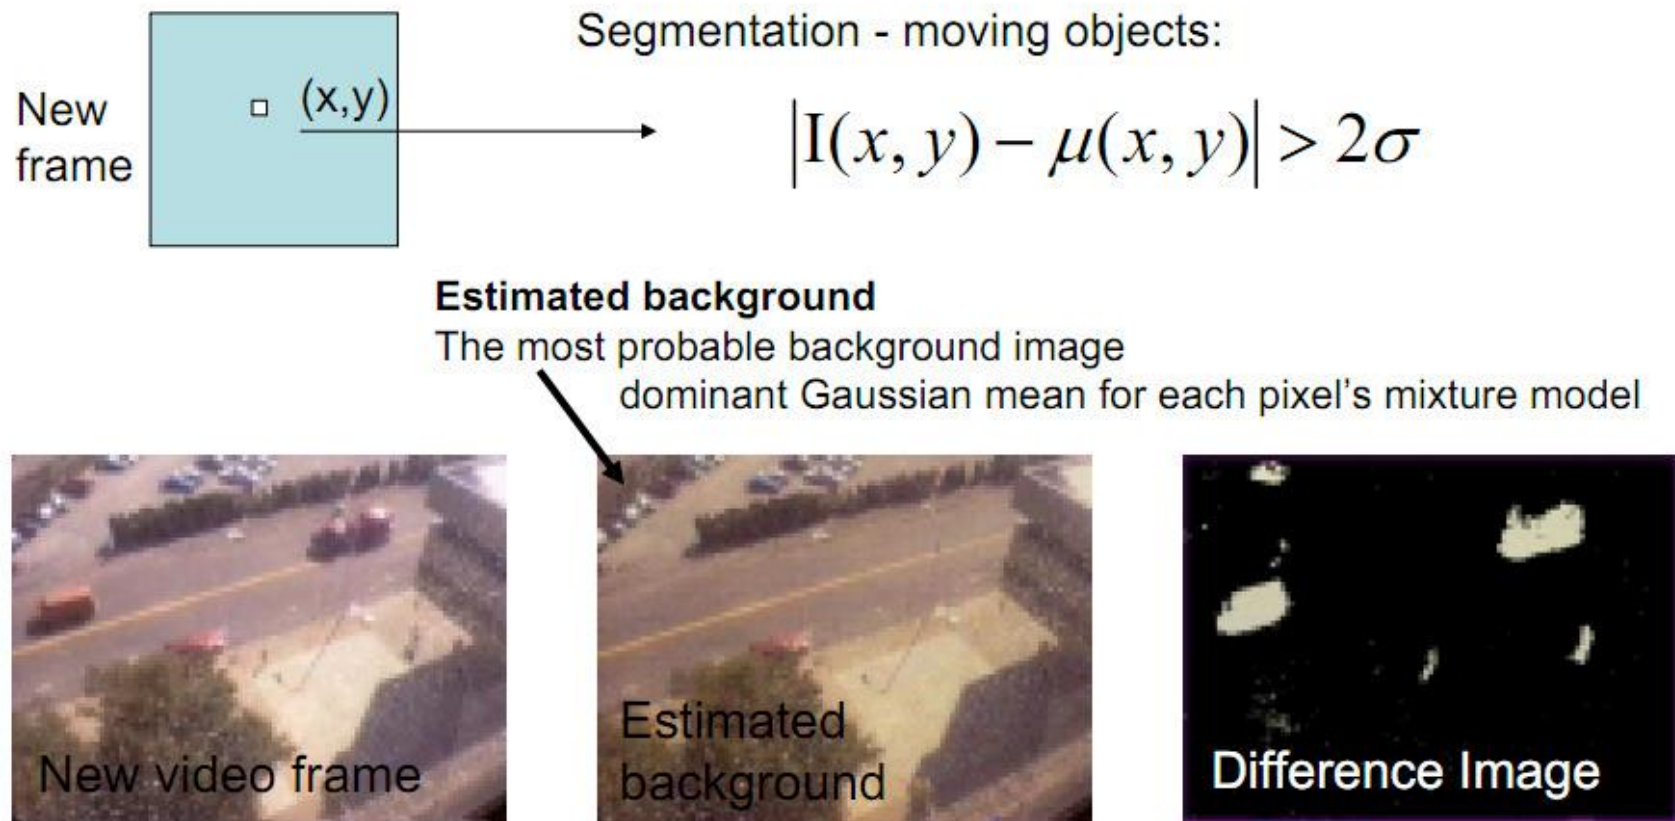
- Examples

|                           |   |   |  |   |   |   |   |
|---------------------------|---|---|--|---|---|---|---|
| Test Image                |    |    |    |    |    |    |    |
|                           | Chair moved   | Light gradually brightened  | Light just switched on   | Tree Waving   | Foreground covers monitor pattern   | No clean background training  | Interior motion undetectable  |
| Ideal Foreground          |    |    |    |    |    |    |    |
| Adjacent Frame Difference |   |   |   |   |   |   |   |
| Mean & Threshold          |  |  |  |  |  |  |  |

- Adaptive Mixture of Gaussians



# Background mixture model - example





- Motion
  - Motion analysis
  - Optical Flow
- Background Subtraction
- Tracking
  - Object Tracking
  - Template Matching
  - Detection vs Tracking





- Object tracking is a **crucial issue** in computer vision, especially for the applications where the environment is in continuous changing:
  - Robot Vision - mobile robot navigation, applications that must deal with unstable grasps, . . .
  - Surveillance
  - Traffic applications
  - Human motion capture

# Tracking: some applications



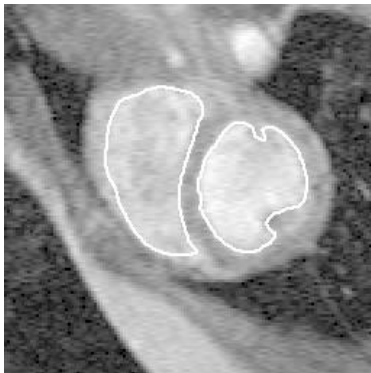
Body pose tracking,  
activity recognition



Censusing a bat  
population



Video-based  
interfaces



Medical apps

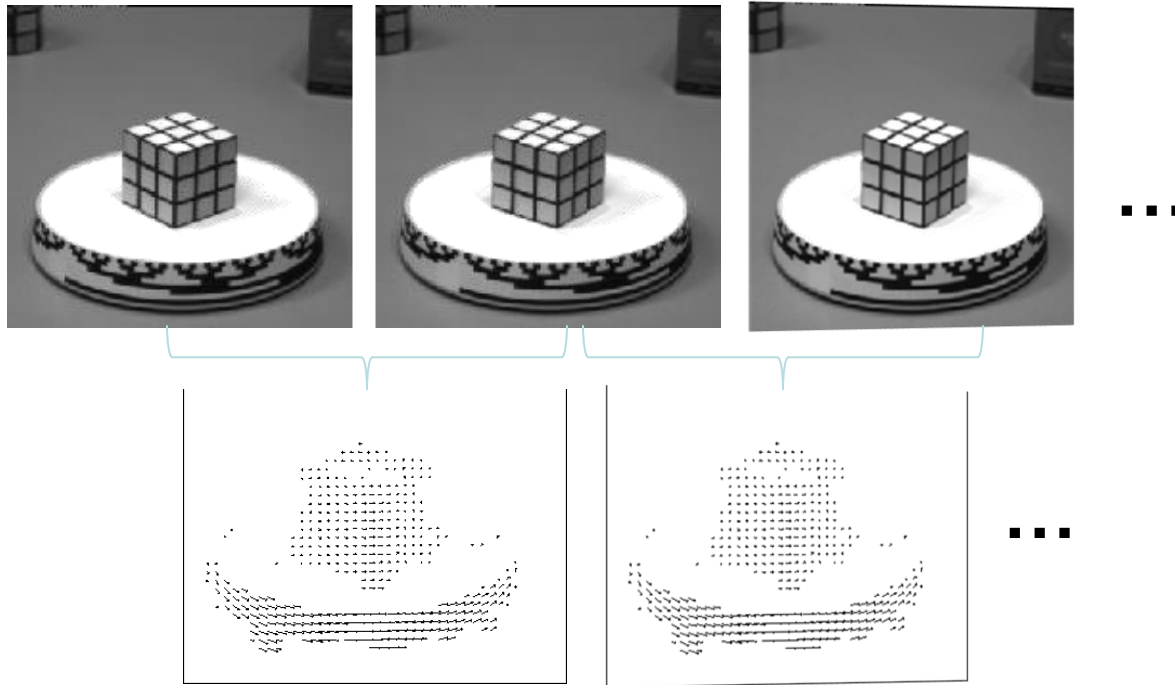


Surveillance

# Optical flow for tracking?



If we have more than just a pair of frames, we could compute **optical flow** from one to the other next

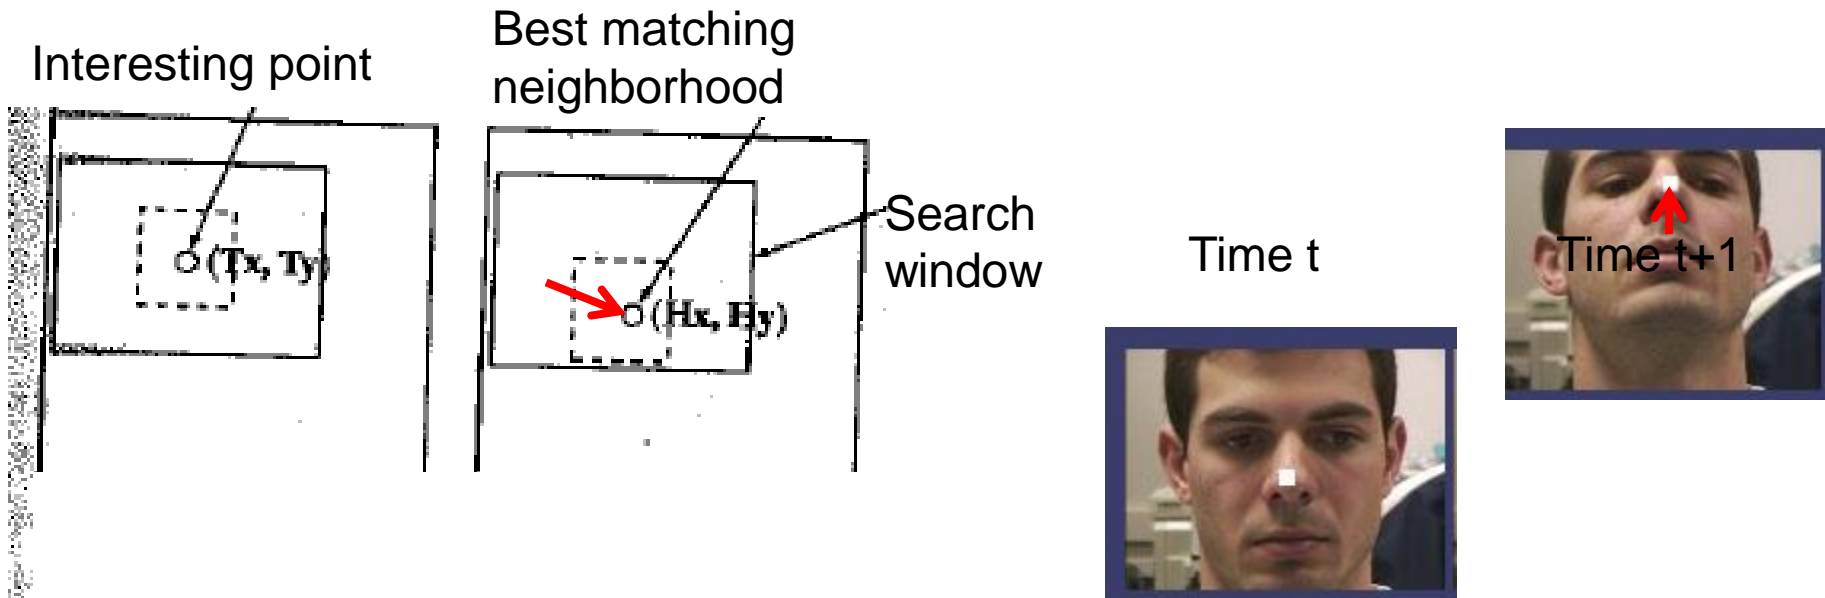


But flow only reliable for small motions, and we may have occlusions, texture less regions that yield bad estimates anyway...

# Feature-based matching for motion



Block matching or template matching (see previous class)



Search window is centered at the point where we last saw the feature, in image I1.

Best match = position where we have the highest normalized cross-correlation value.

# Example: A Camera Mouse



- Video interface: use feature tracking as mouse replacement



- User **clicks** on the feature to be tracked
- Take the **15x15 pixel square** of the feature
- In the next image do a **search** to find the 15x15 **region** with the **highest correlation**
- Move the **mouse pointer** accordingly
- Repeat in the background every **1/30th** of a second



t=1



t=2

...



t=20




t=21



Detection: We **detect** the object **independently** in each frame and can record its position over time, e.g., based on blob's centroid or detection window coordinates





Tracking with *dynamics*: We use image measurements to estimate position of object, but also incorporate position predicted by dynamics, i.e., our expectation of object's motion pattern.



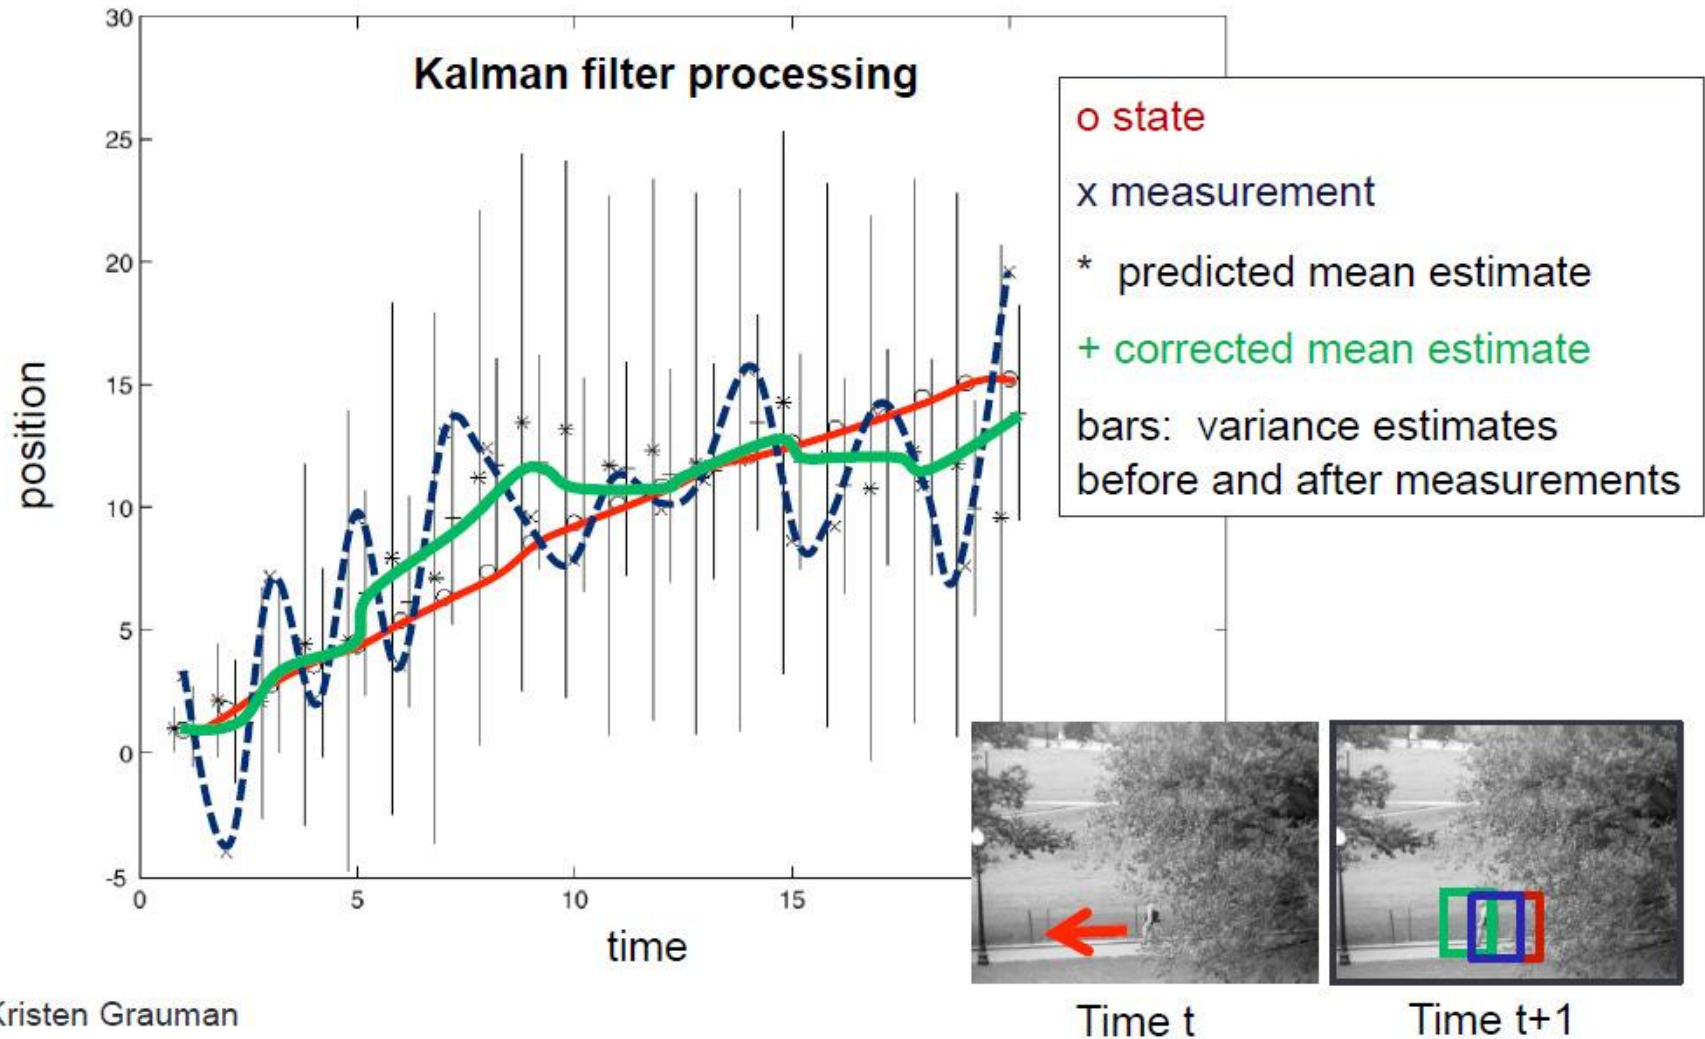


- The **hidden state** consists of the **true parameters** we care about, denoted  $X$ .
- The **measurement** is **our noisy observation** that result from the underlying state, denoted  $Y$ .
- At each time step, **state changes** (from  $X_{t-1}$  o  $X_t$  ) and we get a **new observation**  $Y_t$  .



- Method for **tracking linear** dynamical **models** in Gaussian noise
- The predicted/corrected state distributions are Gaussian
- Only need to **maintain the mean and covariance**

# Tracking: Constant velocity model




Kristen Grauman

- **Initialization**
  - Often done manually
  - Background subtraction, detection can also be used
- **Data association**, multiple tracked objects
  - Occlusions, clutter



- **Deformable** and articulated objects
- Constructing **accurate models** of dynamics - (example: Fitting parameters for a linear dynamics model)
- **Drift** - accumulation of errors over time





- Barron, “Tutorial: Computing 2D and 3D Optical Flow.”, <http://www.tina-vision.net/docs/memos/2004-012.pdf>
- CVonline: Optical Flow
- Fast Image Motion Estimation Demo  
<http://extra.cmis.csiro.au/IA/changs/motion/>
- Part of this lecture is from the courses offered by Prof. Shree Nayar at Columbia University, USA and by Prof. Srinivasa Narasimhan at CMU, USA.